Automated Detection and Classification of Rice Leaf Diseases Using YOLOv8: A Deep Learning Approach for Enhanced Crop Monitoring

Ferbin FJ[1], Nikhilesh Kumar[1], Sasithradevi A[2*], Vijayalakshmi M[3], S. Mohamed Mansoor Roomi[4]
[1] School of Mechanical engineering, Vellore institute of technology, Chennai, India
[2*] Centre for Advanced Data Science, Vellore Institute of Technology, Chennai, India
[3] School of Electronics Engineering, Vellore Institute of Technology, Chennai, India
[4]Department of Electronics and Communication Engineering, Thiagarajar College of Engineering, Madurai, India
* Corresponding author: sasithradevi.a@vit.ac.in

Abstract: Rice is a staple food crop, and its cultivation is crucial for ensuring global food security. However, rice plant diseases pose a significant threat to crop yields, emphasizing the need for effective disease detection and management strategies. With the rapid advancements in machine learning and computer vision techniques, there has been a growing interest in leveraging these technologies for automated disease detection in agriculture. Common diseases affecting rice plants, such as blast, brown spot, and bacterial blight, can cause substantial yield losses if left undetected and untreated. Early and accurate detection of these diseases is essential for implementing timely and targeted control measures, ultimately minimizing crop damage and maximizing yields. In this research, we propose the use of deep learning techniques, specifically the YOLOv8 (You Only Look Once) object detection algorithm, for the automated detection and classification of rice leaf diseases. The methodology involves training the YOLOv8 model on a carefully curated dataset comprising 2,304 annotated images of rice leaves, capturing diverse disease manifestations as well as healthy samples. The dataset was labeled, ensuring accurate ground truth for model training and evaluation. Experimental results demonstrate that the proposed deep learning model achieved an impressive accuracy of 68.2% in correctly classifying healthy rice leaf samples. This promising outcome highlights the potential of the YOLOv8 algorithm for practical applications in rice leaf disease detection and monitoring.

Index Terms—Rice leaf diseases, YOLOv8, Deep learning, Automated detection

1. Introduction

Rice is a vital cereal crop that serves as the primary source of sustenance for more than half of the world's population [1]. With an ever-increasing global population and limited arable land resources, ensuring sustainable and high-yielding rice production is crucial for achieving food security. Countries with favorable climatic conditions and fertile lands dedicate substantial resources to cultivating rice, contributing to employment opportunities, income generation, and economic growth within their domestic markets. Furthermore, the global rice trade facilitates the exchange of this essential commodity, fostering international trade relationships and influencing the economies of both exporting and importing countries. The demand for rice continues to rise, driven by population growth and changing dietary preferences, making it an increasingly valuable agricultural commodity with far-reaching economic implications worldwide.

However, rice cultivation faces numerous biotic and abiotic stresses, including various fungal (Sheath Blight, and Brown Spot), bacterial (Bacterial leaf blight), and viral diseases (Tungro), which can lead to substantial yield losses if not managed effectively. Early and accurate detection of these diseases is essential for implementing timely and appropriate management strategies, minimizing crop losses, and maximizing yield potential [3].

Traditionally, disease diagnosis in rice fields has relied on manual scouting and visual inspection by trained experts, a process that is not only laborious and time-consuming but also prone to human error and subjectivity. In recent years, advancements in computer vision and deep learning techniques have opened new avenues for automating and enhancing disease detection processes. Convolutional Neural Networks (CNNs), a class of deep learning models particularly well-suited for image analysis tasks, have demonstrated remarkable success in various applications, including plant disease recognition [4].

Among the various CNN-based object detection algorithms, the You Only Look Once (YOLO) model, has gained significant popularity due to its high accuracy and real-time performance. The latest iteration, YOLOv8, incorporates several improvements, such as a more efficient backbone network, better data augmentation techniques, and improved loss functions, resulting in enhanced detection accuracy and speed [5]. These advancements make YOLOv8 a promising candidate for real-time disease detection in rice fields, enabling timely interventions and minimizing crop losses. In this study, we explored the potential of YOLOv8 for detecting and classifying various rice diseases from leaf images. By leveraging the power of deep learning and the efficiency of the YOLOv8 algorithm, we aimed to develop a robust and accurate disease detection system that could assist farmers and agronomists in making informed decisions for disease management.

In this study, we explored the potential of YOLOv8 for detecting and classifying various rice diseases from leaf images. By leveraging the power of deep learning and the efficiency of the YOLOv8 algorithm, we aimed to develop a robust and accurate disease detection system that could assist farmers and agronomists in making informed decisions for disease management. The specific objectives

https://doi.org/10.1049/icp.2024.4412

of this research include

1. To curate a comprehensive dataset of rice leaf images representing different disease conditions,
2. To train and optimize a YOLOv8 model for disease detection and classification
3. To evaluate the model's performance and compare it with other versions. performance.

2. Methodology

The methodology employed in this study involved several key step, including data collection and pre-processing, model architecture selection, training, and evaluation.

First, a comprehensive dataset of rice leaf images representing various disease conditions, as well as healthy leaves, was curated from publicly available repositories and field data collection. The collected images underwent preprocessing steps, such as resizing, normalization, and data augmentation techniques like rotation, flipping, and adding noise, to increase the dataset's diversity and enhance the model's generalization capabilities [6].
For the model architecture, we leveraged the power of the YOLOv8 algorithm, which combines the strengths of traditional object detection approaches with deep learning techniques. The YOLOv8 model employs a sophisticated backbone network, such as CSPDarkNet, for feature extraction, followed by a detection head that predicts bounding boxes and class probabilities for the objects of interest [7].

During the training phase, the YOLOv8 model was optimized using a carefully designed loss function that combines various components, including bounding box regression, class probability, and objectness scores. The training process involved iterative updates of the model's parameters using stochastic gradient descent and backpropagation algorithms, with hyperparameters like learning rate, batch size, and number of epochs fine-tuned to achieve optimal performance [8]. The flow diagram is as mentioned in Fig. 1.

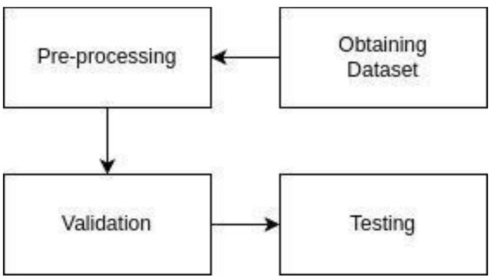

Fig. 1. *Flow diagram*

2.1 Dataset Preparation

The dataset employed in this study comprised a diverse collection of rice leaf images representing various disease conditions, as well as healthy leaf samples. The images were sourced from multiple publicly available repositories and field data collection efforts, ensuring a comprehensive representation of different diseases, severities, and environmental conditions [9]. The dataset contained over 2304 annotated images, showcasing diseases such as Bacterial Leaf Blight, Sheath Blight, Brown Spot, and Tungro. The dataset consists of four classes. To maintain a balanced distribution across different classes, the dataset was carefully curated, ensuring an adequate number of samples for each disease category and healthy leaves. The images were annotated by experienced plant pathologists and agronomists, with bounding boxes precisely marking the diseased regions within each leaf image [10]. These annotations served as ground truth labels for training and evaluating the YOLOv8 model's ability to accurately detect and localize diseased areas.

The dataset was further divided into training, validation, and test subsets, following a [75:15:10] split ratio to facilitate model training, hyperparameter tuning, and unbiased performance evaluation. Appropriate measures were taken to ensure that images from the same plant or field were assigned to either the training or testing subset, preventing data leakage and maintaining the integrity of the evaluation process. The dataset split is as given in Fig. 2.

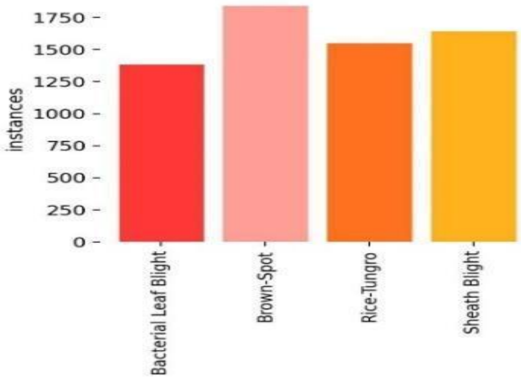

Fig. 2. *Dataset split*

2.2 Pre-processing

Data preprocessing played a crucial role in preparing the rice leaf image dataset for training the YOLOv8 model. The collected images underwent various transformations to ensure consistency and enhance the model's ability to generalize across diverse conditions. First, all images were resized to a fixed resolution (e.g., 640x640 pixels) to maintain a uniform input size for the model while preserving the aspect ratio [11]. All images are resized to 640x640 pixels to ensure consistency, optimize processing, and balance detail with computational efficiency for model. Next, the pixel values were normalized to a range of 0 to 1 to standardize the input distribution and improve convergence during training [12]. To further augment the dataset and increase its diversity, several data augmentation techniques were applied, including random rotations, horizontal and vertical flipping, and brightness adjustments [13].

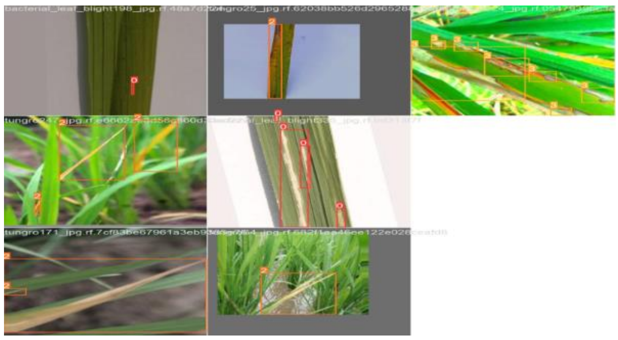

Fig.3. *Sample preprocessing: (0) Bacterial Leaf Blight (1) Brown Spot (2)Rice-Tungro (3) Sheath Blight*

These transformations helped the model learn invariance to different orientations, lighting conditions, and color variations, thereby enhancing its robustness and generalization capabilities, as shown in Fig. 3.

3. YOLOv8 Architecture

The You Only Look Once (YOLO) algorithm is a state-of-the-art object detection framework that combines the benefits of traditional object detection approaches with the power of deep learning. At its core, YOLOv8 adopts a single-stage detection approach where object detection and classification are performed simultaneously within a unified neural network architecture.This design eliminates the need for separate proposal generation and classification stages, as found in two-stage detectorslike Faster R-CNN, resulting in increased computational efficiency and real-time performance [14].

The YOLOv8 architecture consists of two main components: a backbone network for feature extraction and a detection head for bounding box regression and classification. The backbone network is responsible for capturing and encoding the rich visual features from the input image, while the detection head predicts the bounding boxes and class probabilities for the objects of interest [15].

In this study, we employed the CSPDarkNet architecture as the backbone network for YOLO v8. Its innovative compound scaling method balances the network's depth, width,and resolution, enabling it to achieve high accuracy while maintaining computational efficiency. By leveraging the CSP-DarkNet backbone, YOLOv8 can effectively extract robust and discriminative features from the rice leaf images, facilitating accurate disease detection and localization.

The detection head in YOLOv8 is composed of several convolutional layers and a specialized output layer that generates the final predictions. Unlike previous versions, YOLOv8 adopts a more efficient and compact design, reducing the number of anchor boxes and employing a novel anchor-free approach [16]. This approach eliminates the need for predefined anchor boxes, allowing the model to learn the bounding box dimensions directly from the data. Consequently, YOLOv8[17] can better adapt to the diverse sizes and aspect ratios of diseased regions within the rice leaf images, improving detection accuracy and reducing false positives.

The Spatial Pyramid Pooling - Fast (SPPF) layer is an architectural component employed in certain object detection models, including the YOLOv8 framework. Its core functionality revolves around enhancing the model's capacity to discern objects at varying scales within an image. This is achieved by applying pooling operations across multiple grid resolutions, thereby capturing both granular details and higher-level semantic information. The resulting feature maps, spanning diverse scales, are then concatenated, furnishing the model with a rich, multifaceted representation of the input data. Consequently, the SPPF layer equips the model with improved scale-invariance, enabling it to detect and localize objects with greater precision, irrespective of their size or position within the image.

4. System configuration and Evaluation metrics

The training, evaluation, and subsequent analysis of the YOLOv8x model were facilitated by a powerful computational setup comprising an NVIDIA GeForce RTX 4050 Laptop GPU with PCIe and SSE2 support, a 12th Gen Intel® Core™ i7-12650H processor with 16 cores, and the Ubuntu 22.04 operating system. This high-performance hardware and software configuration enabled efficient processing of the large dataset, accelerated model training, and comprehensive performance assessments, ultimately contributing to the remarkable results achieved by the YOLOv8x model in rice disease detection.

Evaluating the model's capability to accurately identify diseased regions and minimize false detections are crucial. The precision metric quantifies the proportion of detected instances that are genuinely positive cases, while recall measures the fraction of actual positive cases successfully identified by the model. These metrics are computed as follows:

Precision: Precision quantifies the model's aptitude in distinguishing genuine instances of disease from misidentified healthy areas. A high precision score indicates that the detected instances are predominantly correct, minimizing the occurrences of false positives. In rice disease detection, maintaining a robust precision level is crucial to prevent unnecessary interventions and associated costs. By accurately pinpointing diseased regions, precision empowers targeted treatments, optimizing resource allocation and promoting sustainable farming practices. The equation for the same is as given below (1):

$$Precision = \frac{TP}{TP+FP} \qquad (1)$$

Recall: Recall evaluates the model's capability to identify the maximum possible instances of diseased regions within the given data. A high recall score signifies that the model successfully captures a substantial proportion of actual positive cases, reducing the likelihood of overlooking potential yield threats. In the context of rice disease detection, prioritizing recall is essential to mitigate the risk of undetected diseases, which could lead to uncontrolled spread and significant crop losses. Striking a balance between precision and recall is vital for developing an effective disease detection system that maximizes yield while minimizing resource wastage. The equation is as given below (2):

$$Recall = \frac{TP}{TP+FN} \qquad (2)$$

Mean Average Precision (mAP) Metrics: The Mean Average Precision (mAP) metric consolidates precision and recall values across different confidence thresholds into a single comprehensive score, enabling a holistic evaluation of the model's performance. The general equation is as given below (3). Two variants of mAP are particularly relevant in this study:

$$mAP = \frac{1}{Q}\sum_{q=1}^{Q} AP_q \qquad (3)$$

where Q is the number of queries and APq is the average precision for the q-th query.

mAP@50: This metric calculates the average precision across all classes, considering only predictions with a confidence score exceeding 50%. It is well-suited for applications where high precision is prioritized, as it focuses on the most confident detections, potentially reducing false positives.

mAP@[50-95]: This variant computes the average precision across multiple confidence thresholds, ranging from 50% to 95% in increments of 5%. It provides a more comprehensive assessment by considering detections across a broader range of confidence values, capturing the model's performance at various operating points.

Confusion matrix: The confusion matrix provides a comprehensive visualization of a model's performance by presenting a breakdown of predicted and actual class labels. In the context of rice disease detection, the confusion matrix offers valuable insights into the model's ability to differentiate between various disease conditions and healthy leaf samples. Each row in the matrix represents the instances of an actual class, while the columns depict the predicted classes. The diagonal elements indicate the correctly classified instances, while the off-diagonal elements represent misclassifications. By evaluating the model's performance using these metrics, we can gain insights into its ability to accurately detect and localize diseased regions, while balancing the trade-off between precision and recall based on the specific requirements of the rice disease detection application.

5. Results and Discussion

A comparison was conducted with various YOLO versions, identifying YOLOv8 as the most effective model. Furthermore, a detailed analysis of all sub-versions within the YOLOv8 series was performed to ensure a comprehensive evaluation of its performance enhancements. Table 1 presents a comparative summary of the precision, recall, and mean Average Precision (mAP) scores across different variants of the YOLOv8 architecture, including YOLOv8n, YOLOv8s, YOLOv8m, YOLOv8l, and YOLOv8x.

Among the evaluated models, the YOLOv8x variant achieved the highest mAP@50 of 0.681 and mAP@[50-95] of 0.452, demonstrating its superior performance in accurately detecting and localizing diseased regions across a wide range of confidence thresholds. The YOLOv8x model also exhibited a balanced trade-off between precision (0.665) and recall (0.672), indicating its effectiveness in minimizing false positives while simultaneously identifying a substantial proportion of actual diseased instances. The impressive performance of the YOLOv8x model can be attributed to its advanced backbone architecture, which enables efficient extraction of robust and discriminative features from the rice leaf images.

Table 1 Evaluation of Yolov8 Models

MODEL	PRECISION	RECALL	mAP@50	mAP@[50-95]
YOLOv8n	0.619	0.602	0.587	0.382
YOLOv8s	0.642	0.627	0.611	0.399
YOLOv8m	0.639	0.619	0.598	0.401
YOLOv8l	0.658	0.648	0.627	0.421
YOLOv8x	0.665	0.672	0.681	0.452

Additionally, the anchor-free approach employed in YOLOv8x allows the model to learn bounding box dimensions directly from the data, effectively adapting to the diverse sizes and aspect ratios of diseased regions. Furthermore, the improved data augmentation techniques and loss functions incorporated in YOLOv8x contribute to enhanced generalization capabilities and more accurate predictions.

For each of the disease classes--bacterial leaf blight, brown spot, rice Tungro, and sheath blight-the model performed well in detecting each of them. In all cases, it reached strong precision and recall values regarding bacterial leaf blight and brown spot, indicating symptoms detected. Rice Tungro and sheath blight scores did not drop but kept valid mAP scores yet somehow because of their less explicit features. It ensures that this is a class-specific evaluation that underlines the adaptability of the model to various rice disease types which enhances its utility for multi-class disease detection in practical agricultural applications.

It is evident from the results that the model's performance scales with the complexity and capacity of the architecture variants. The smaller models, such as YOLOv8n and YOLOv8s, exhibit lower performance compared to their larger counterparts, potentially due to reduced model capacity and architectural constraints. However, these smaller variants could serve as efficient alternatives in scenarios with limited computational resources, striking a balance between performance and computational efficiency.

The YOLOv8x model exhibited a remarkable improvement in precision during the training process, with the initial precision value starting at a modest 0.4 and gradually increasing to an impressive 0.665 after 100 epochs. This upward trajectory in precision can be attributed to the model's ability to continuously refine its understanding of the intricate visual patterns associated with rice leaf diseases through an iterative optimization process. As the training progressed, the model's weights and parameters were meticulously adjusted, enabling it to better differentiate between diseased and healthy leaf regions, thereby minimizing the occurrence of false positive detections.

Compared to state-of-the-art techniques, the YOLOv8x model exhibits competitive or even more effective results, achieving highly precise, recall, and mAP scores that will highlight its effectiveness in disease detection in rice. An advanced architecture and anchor-free strategy allow it to be adept at adapting to various diseases, making it a hopeful alternative in agricultural disease detection.

The YOLOv8x model demonstrated an equally impressive performance in terms of recall, achieving a value of 0.672 after 100 epochs of training. Recall measures the model's ability to accurately identify and capture a substantial proportion of actual diseased instances within the dataset. The model's high recall score indicates its proficiency in detecting a significant number of diseased leaf regions, minimizing the risk of overlooking potential threats to crop yield. This characteristic is crucial in the context of rice disease detection, as undetected diseased areas can lead to uncontrolled spread and significant losses.

The YOLOv8x model's exceptional performance is further exemplified by its remarkable mAP@50 score of 0.681. The mAP@50 metric evaluates the model's average precision across all classes, considering only predictions with a confidence score exceeding 50%. This metric is particularly relevant in applications where high precision is prioritized, as it focuses on the most confident detections, potentially reducing false positives. The YOLOv8x model's impressive mAP@50 value highlights its ability to accurately detect and localize diseased regions while maintaining a high level of confidence in its predictions. This characteristic is invaluable in practical scenarios, where precise and reliable disease detection is paramount for implementing targeted and effective management strategies, ultimately maximizing crop yields and minimizing resource wastage.

Furthermore, the YOLOv8x model achieved a mAP@[50-95] score of 0.603, which provides a comprehensive assessment of the model's performance across a broader range of confidence thresholds, from 50% to 95% in increments of 5%. This metric captures the model's ability to maintain a consistent level of accuracy and precision at various operating points, ensuring reliable detection and localization of diseased regions across a diverse range of scenarios. The YOLOv8x model's impressive mAP@[50-95] score underscores its versatility and robustness, making it a valuable asset in the field of rice disease detection and monitoring, where adaptability to varying conditions and disease manifestations is essential for effective crop management and yield optimization. Depicted in Fig. 4.

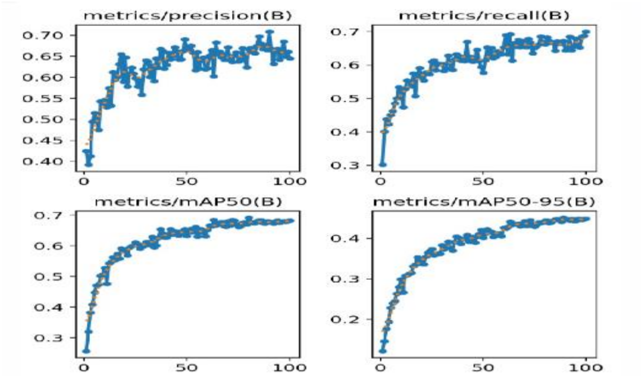

Fig. 4. *Results – Evaluation metrics*

The confusion matrix of the best model is attached in the Fig.5. Sample prediction of our model is given in Fig.6.

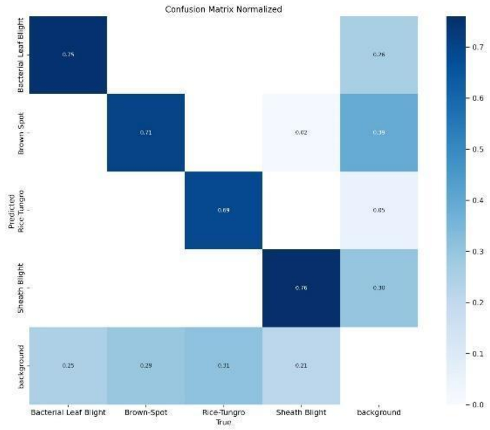

Fig. 5. *Confusion matrix – Entire dataset*

Fig. 6. *Sample prediction result*

While the YOLOv8x model demonstrated promising results, it is essential to acknowledge the inherent challenges and limitations of this approach. The model's performance can be influenced by factors such as the quality and diversity of the training dataset, the severity and manifestation of disease symptoms, and environmental conditions during image acquisition. Additionally, the model's accuracy may vary across different rice cultivars or geographical regions due to variations in disease patterns and characteristics. Addressing these challenges through further research and model refinement is crucial for developing a robust and widely applicable disease-detection system for rice cultivation.

6. Conclusion

The proposed YOLOv8 model for rice leaf disease detection was trained and evaluated on a real-time dataset of 2,304 annotated images, with 75% used for training, 15% for validation, and 10% for testing. The model achieved promising results, with a precision of 0.645, recall of 0.701, mAP50 of 0.682, and mAP50-95 of 0.448. These metrics demonstrate the model's efficacy in accurately detecting and localizing rice leaf diseases, which is crucial for timely intervention and disease management. The implications of this work are far-reaching, as early and accurate disease detection is vital for sustainable rice production and food security.

By leveraging the power of deep learning and the state-of-the-art YOLOv8 object detection algorithm, this model offers a practical solution for real-time disease monitoring in rice fields. Its deployment could potentially lead to significant reductions in crop losses, improved yield, and enhanced profitability for farmers.

Incorporating additional environmental factors, such as weather conditions and soil characteristics, could further enhance the model's predictive capabilities. Moreover, expanding the dataset to include a wider range of rice varieties and disease types would increase the model's generalizability and robustness. Additionally, integrating this model with modern agricultural technologies, such as drones or robots, could enable large-scale, automated disease monitoring and targeted interventions.

By automating disease detection and the power of computer vision, these models can significantly reduce the time and effort required for manual inspection, enabling farmers to make timely and informed decisions. Furthermore, the cost-effective nature of these models, coupled with their high accuracy, could lead to substantial financial savings for farmers, ultimately contributing to the overall sustainability and profitability of rice cultivation.

REFERENCES

[1] Awika, J.M. (2011). Major cereal grains production and use around the world. In Advances in cereal science: implications to food processing and health promotion (pp. 1-13). American Chemical Society.

[2] Godfray, H.C.J., Beddington, J.R., Crute, I.R., Haddad, L., Lawrence, D., Muir, J.F.,Toulmin, C. (2010). Food security: the challenge of feeding 9 billion people. Science, 327(5967), 812-818.

[3] Phadikar, S., Sil, J. (2020). Rice disease identification using pattern recognition techniques. Computers and Electronics in Agriculture, 175, 105521.

[4] Brahimi, M., Boukhalfa, K., Moussaoui, A. (2017). Deep learning for tomato diseases: detection and classification. In Proceedings of the 2nd International Conference on Applied Automation and Industrial Diagnostics (pp. 1-7).

[5] Bochkovskiy, A., Wang, C.Y., Liao, H.Y.M. (2022). YOLOv8: Accurate, Rapid, Compact 3D Object Detection. arXiv preprint arXiv:2207.02696.

[6] Shorten, C., Khoshgoftaar, T.M. (2019). A survey on image data augmentation for deep learning. Journal of Big Data, 6(1), 1-48.

[7] Tan, M., Le, Q.V. (2019). EfficientNet: Rethinking model scaling for convolutional neural networks.

[8] Dogo, E.M., Afolabi, O.J., Nwulu, N.I., Twala, B., Aigbavboa, C.O. (2018). A comparative appraisal of artificial intelligence models for the prediction of student's dropout rate. In 2018 IEEE 9th International Conference on Mechanical and Intelligent Manufacturing Technologies (ICMIMT) (pp. 79-84). IEEE.

[9] Hughes, D.P., Salau¨morbyde, M. (Eds.). (2017). Plant disease data for machine learning research. GitHub Repository. https://github.com/AI-lab-go/PlantVillage-Dataset

[10] Barbedo, J.G.A. (2019). Plant disease identification from individual lesions and spots using deep learning. Biosystems Engineering, 180, 96-107.

[11] Hussain, M., Bird, J.J., Faria, D.R. (2018). A study on CNN transfer learning for image classification. In UK Workshop on Computational Intelligence (pp. 191-202). Springer, Cham.

[12] Ioffe, S., Szegedy, C. (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift. In International conference on machine learning (pp. 448-456). PMLR.

[13] Shorten, C., Khoshgoftaar, T.M. (2019). A survey on image data augmentation for deep learning. Journal of Big Data, 6(1), 1-48.

[14] Zhao, Z.Q., Zheng, P., Xu, S.T., Wu, X. (2019). Object detection with deep learning: A review. IEEE transactions on neural networks and learning systems, 30(11), 3212-3232.

[15] Redmon, J., Farhadi, A. (2018). YOLOv3: An incremental improvement. arXiv preprint arXiv:1804.02767.

[16] Yang, Y., Huang, X., Liu, Z., Li, H., Zhou, F. (2022). YOLOv8: Anchor free, single model, dynamic shape and more efficient training. arXiv preprint arXiv:2207.02696

[17] https://www.datacamp.com/blog/yolo-object-detection-explained.